



seL4 API: Theory

14 August 2015

Adrian Danis



Australian Government



NSW
GOVERNMENT | Trade & Investment



Queensland
Government



seL4 API: Key Concepts



- **Kernel Object**
 - in-kernel datastruct, only directly accessible by kernel
- **Capability**
 - reference to a kernel object
 - allows holder to invoke functions on the objects
 - i.e. ask kernel to do something with the object
 - holder: thread invoking the cap
- **Low-level interface for key activities**
 - create kernel objects
 - create and manage caps in a CSpace
 - create and manage VSpace
 - create and manage threads
 - communicate between threads

- **Kernel-maintained**
 - user-level cannot directly access or manipulate a capability
 - capability is stored in CSpace
 - pass CSpace address of cap in system calls
- **Datatypes**
 - `seL4_CPtr`
 - index into current thread's cspace root (cnode)
 - this can be tricky....

- CNode Object

- consists of slots in which capabilities are stored
- can also store CNode caps in slots, creates hierarchical CSpace structure

- API

- insert cap:

- indirectly: through untyped retype
- seL4_CNode_Copy
- seL4_CNode_Mint
- seL4_CNode_Move

- remove cap

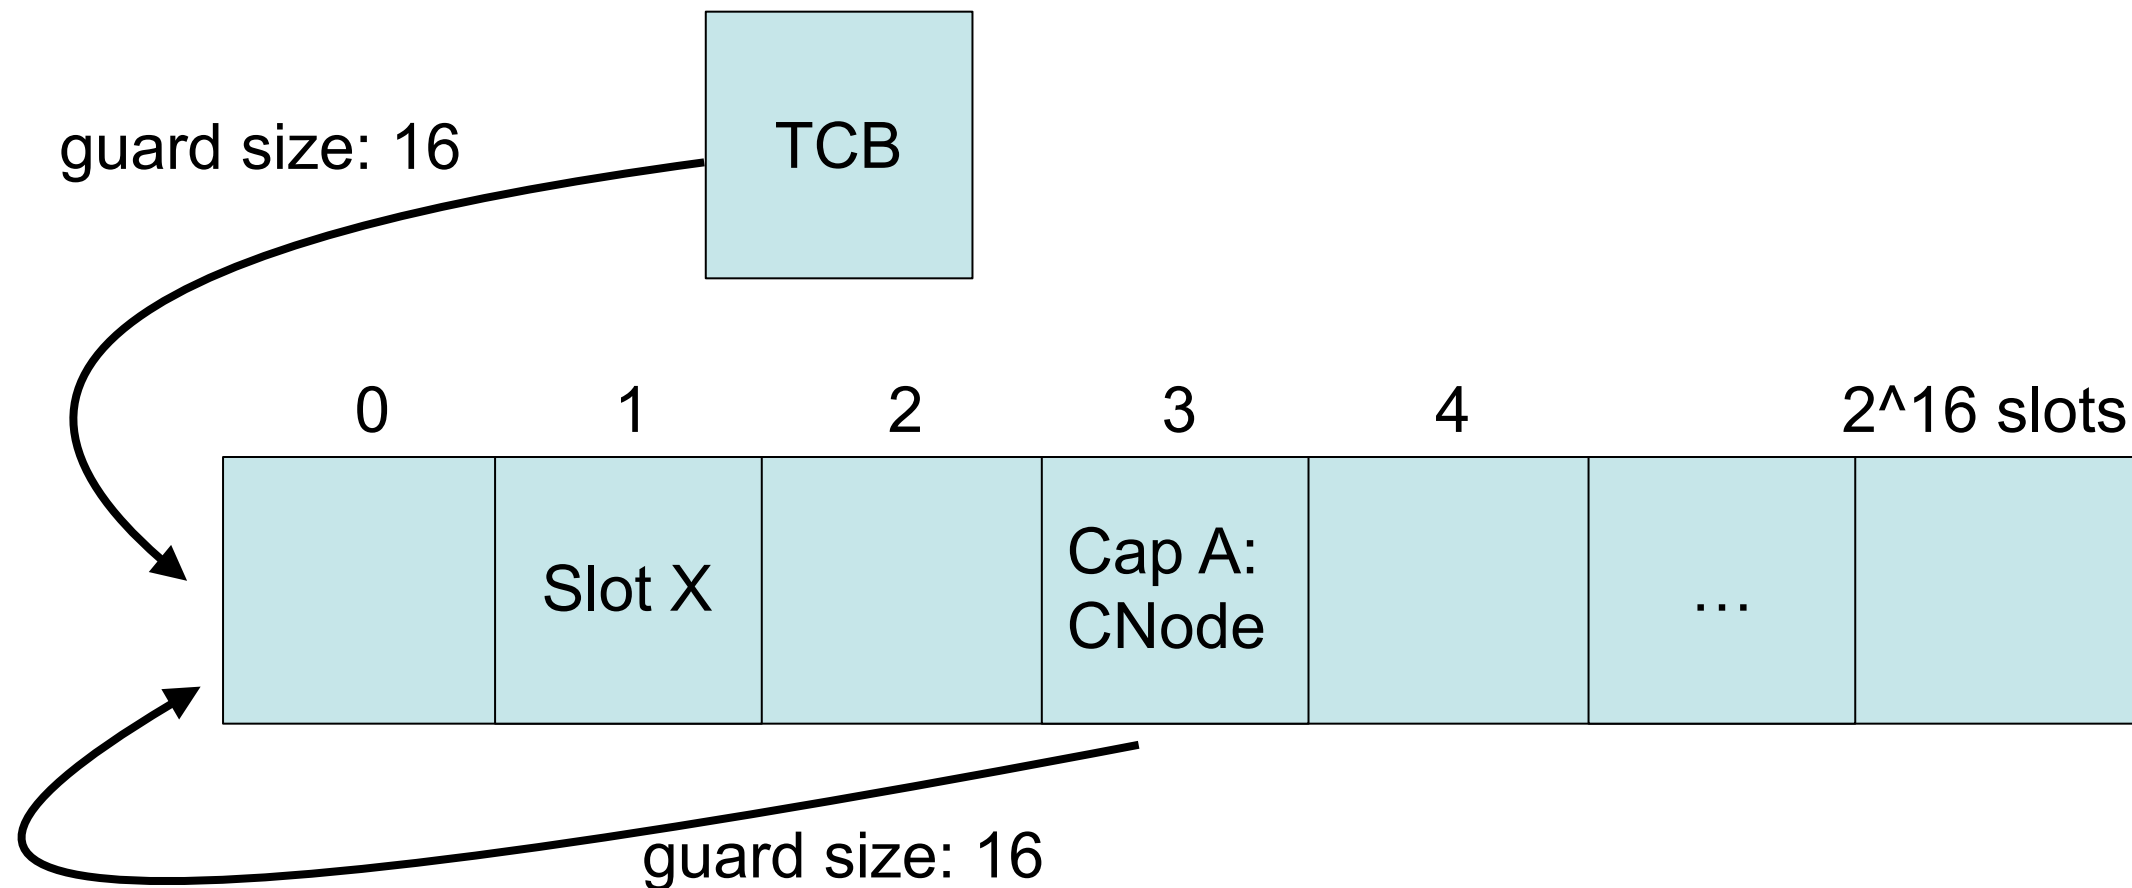
- seL4_CNode_Revoke
- seL4_CNode_Recycle
- seL4_CNode_Delete

CSpace: Addressing



- CSpace structure
 - hierarchy of CNodes
 - CNode has: guard, size (radix)
- CSpace address:
 - refers to a specific slot in the CSpace
- Multi-part address
 - 1. root cnode (seL4_CPtr, relative to TCB cnode, 32-bit)
 - 2. index (into cnode we've derived)
 - 3. depth (bits, amount of index to use)
 - keep resolving index, following CNodes, until
 - no more CNode to follow
 - no more index to resolve
 - 4. slot (if the above refers to a CNode, then a slot in that CNode)

CSpace Addressing Example



Slot X = root: 0x00000003; index 1; depth 32
0x00000003 - guard = 0x0003 = Cap A CNode
index 0x00000001 at depth 32 = Slot X

- **Untyped Memory Object**
 - region of (RAM) memory
 - must be retyped to another object to use it
 - results in a nested tree from a root untyped to other objects:
- **Retyping**
 - kernel uses part of untyped's memory region to store a new kernel object
 - can only create an object if you have a cap to a big enough untyped object
 - retype provides user with cap to the new object
- **seL4_Untyped_Retype**
 - `seL4_Untyped_Retype(seL4_Untyped service, int type, int size_bits, seL4_CNode root, int node_index, int node_depth, int node_offset, int num_objects)`

CSPACE Construction Example

guard size: 8

TCB

2^{16} size

Cap A

Slot X

guard
size: 8

untyped

Cap C

Cap B

Slot Y

guard size: 8


1. Create new CNode with cap in slot Y
2. Copy CNode cap in slot Y to slot X
3. Copy Cap C to a slot in the new CNode

CSpace Construction Example


guard size: 8

TCB

2^{16} size



guard
size: 8



guard size: 8

CSpace Construction Example

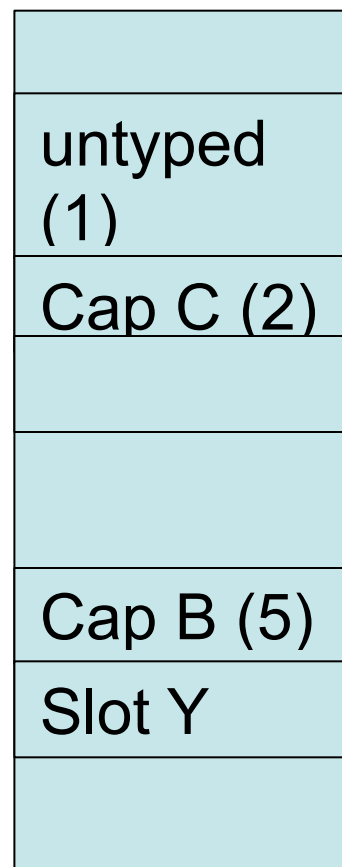
guard size: 8

TCB

2^{16} size



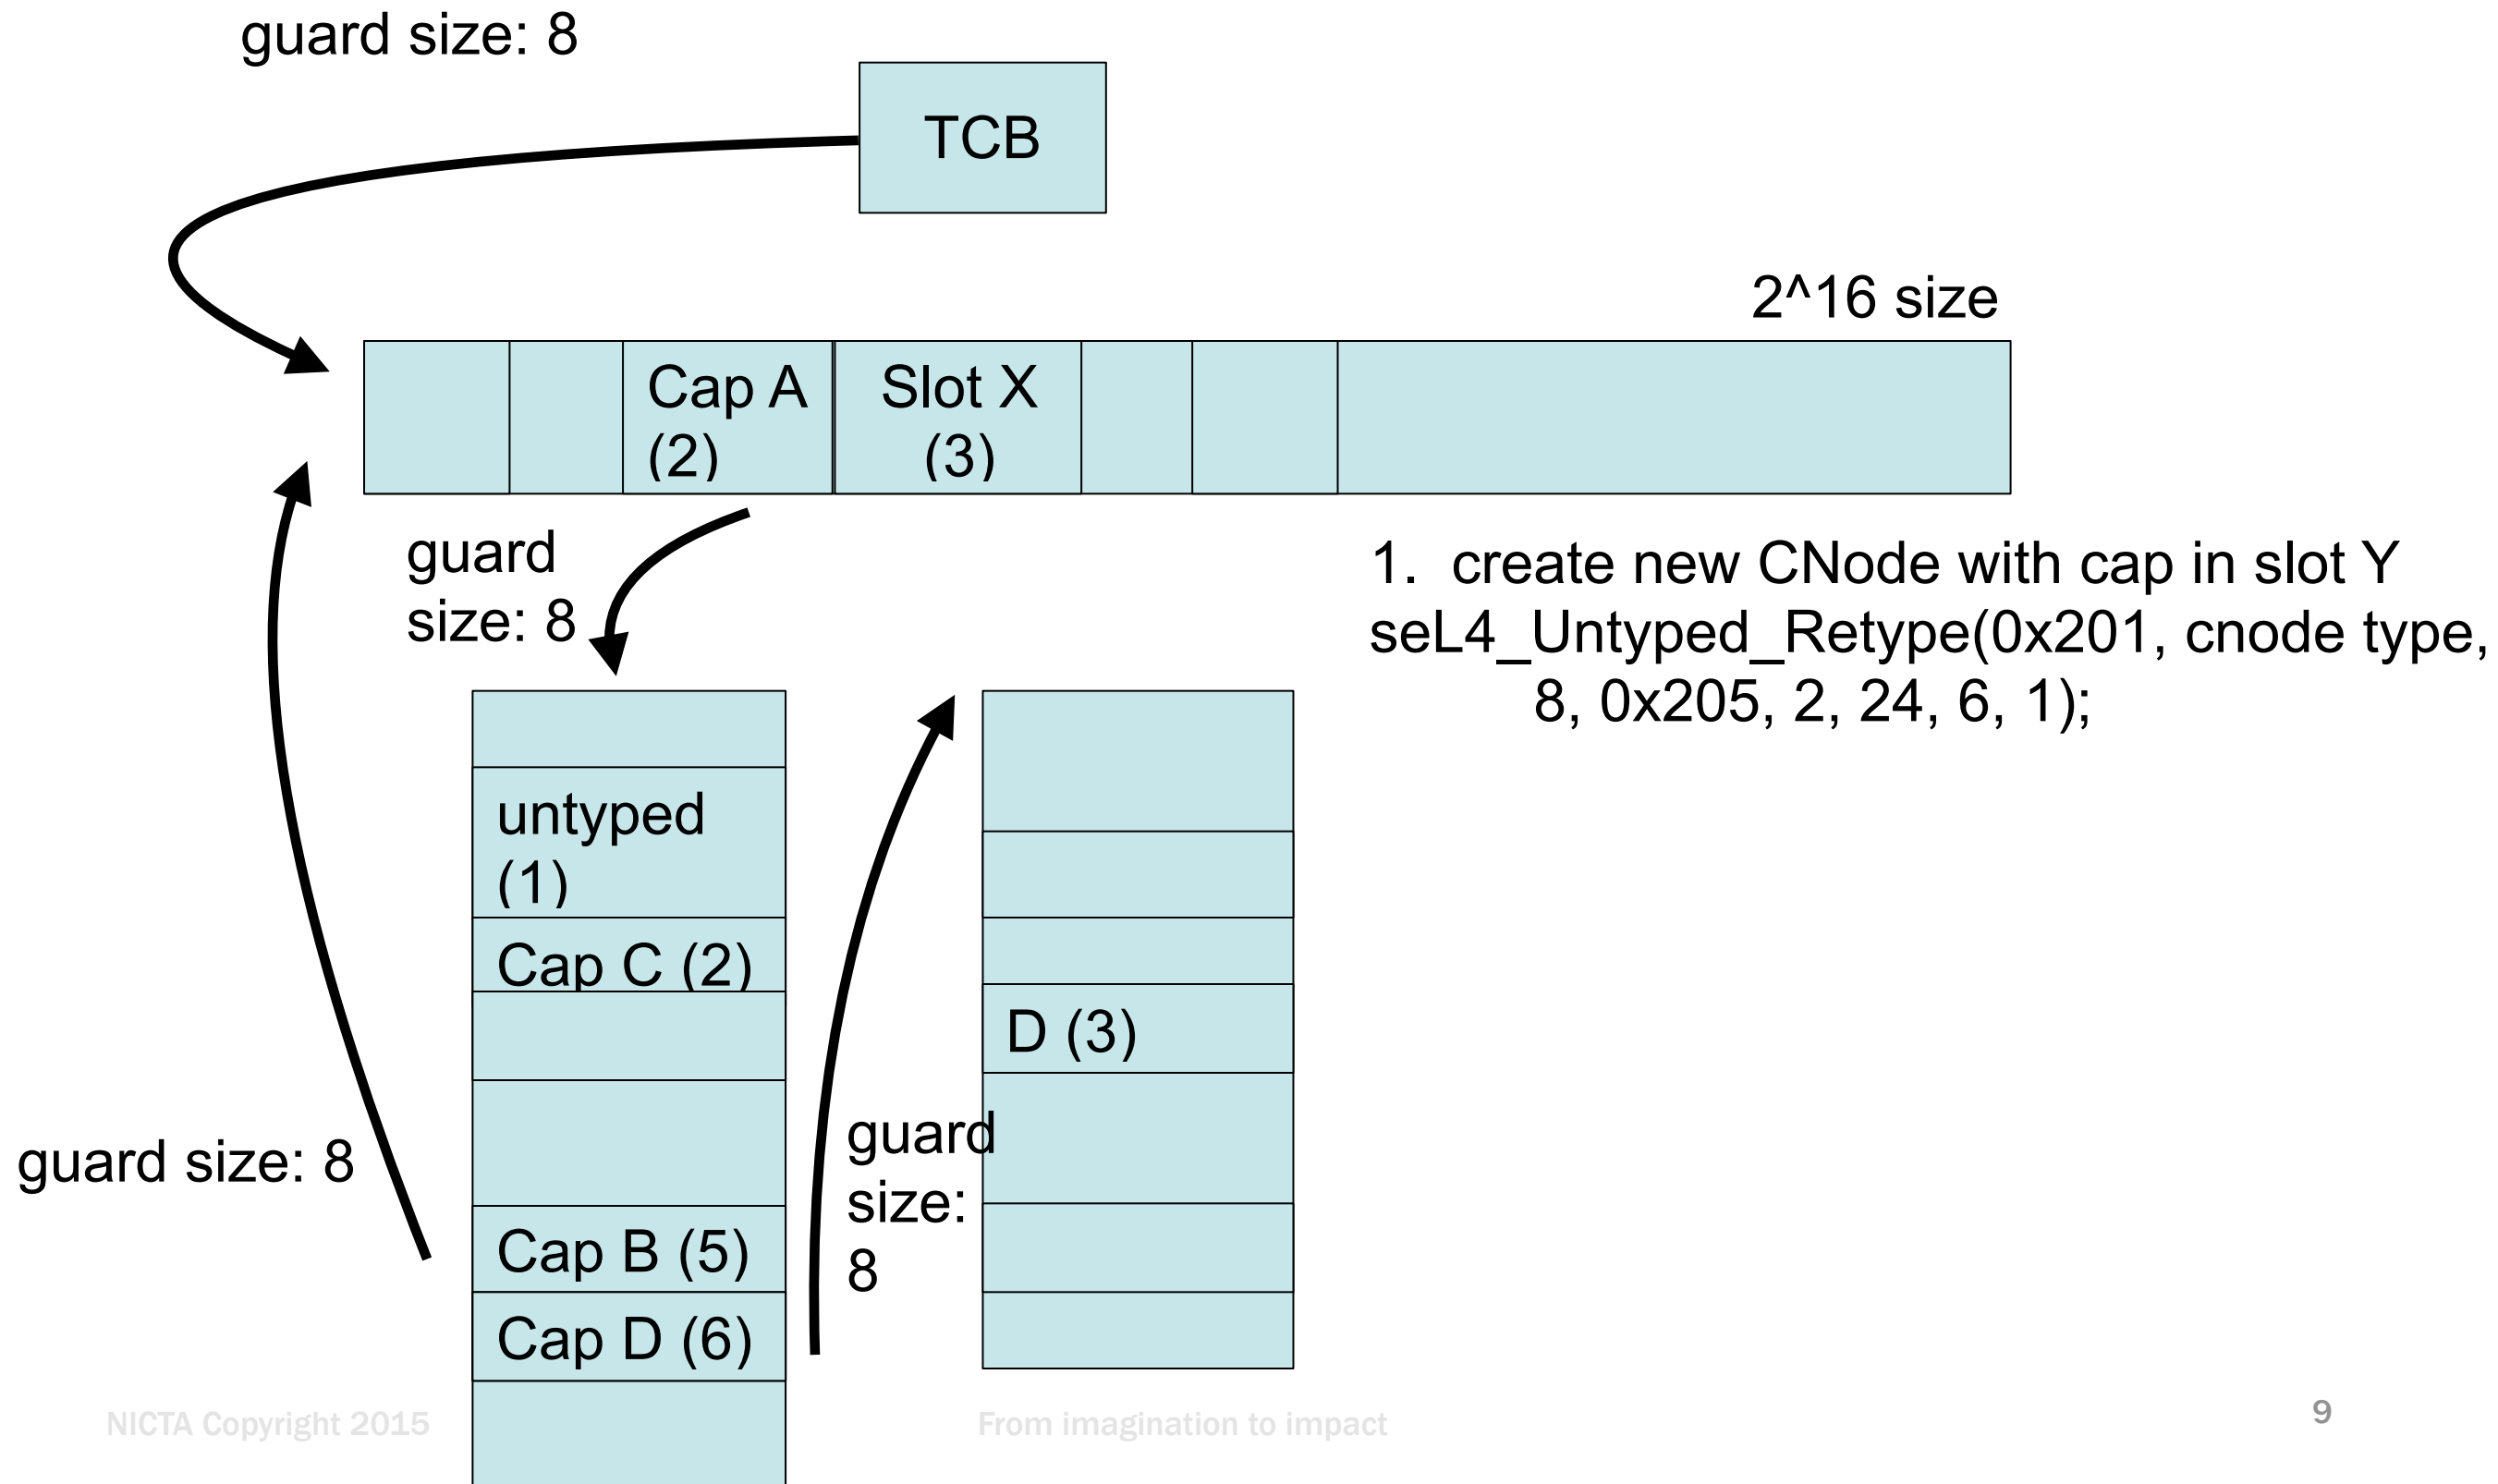
guard size: 8



1. create new CNode with cap in slot Y
`seL4_Untyped_Retype(0x201, cnode type,
8, 0x205, 2, 24, 6, 1);`

guard size: 8

CSpace Construction Example



CSpace Construction Example

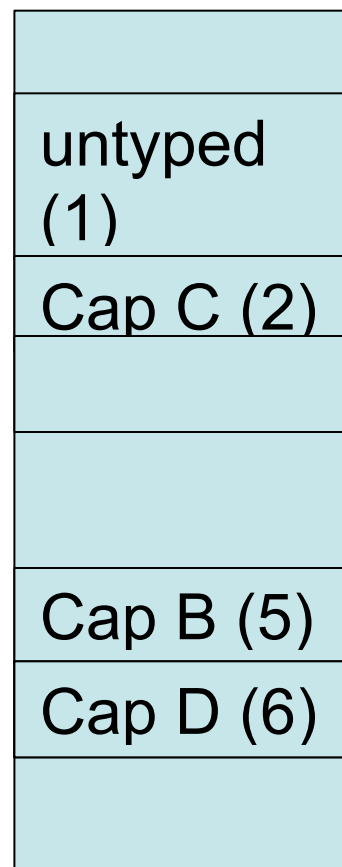
guard size: 8

TCB


2^{16} size



guard size: 8



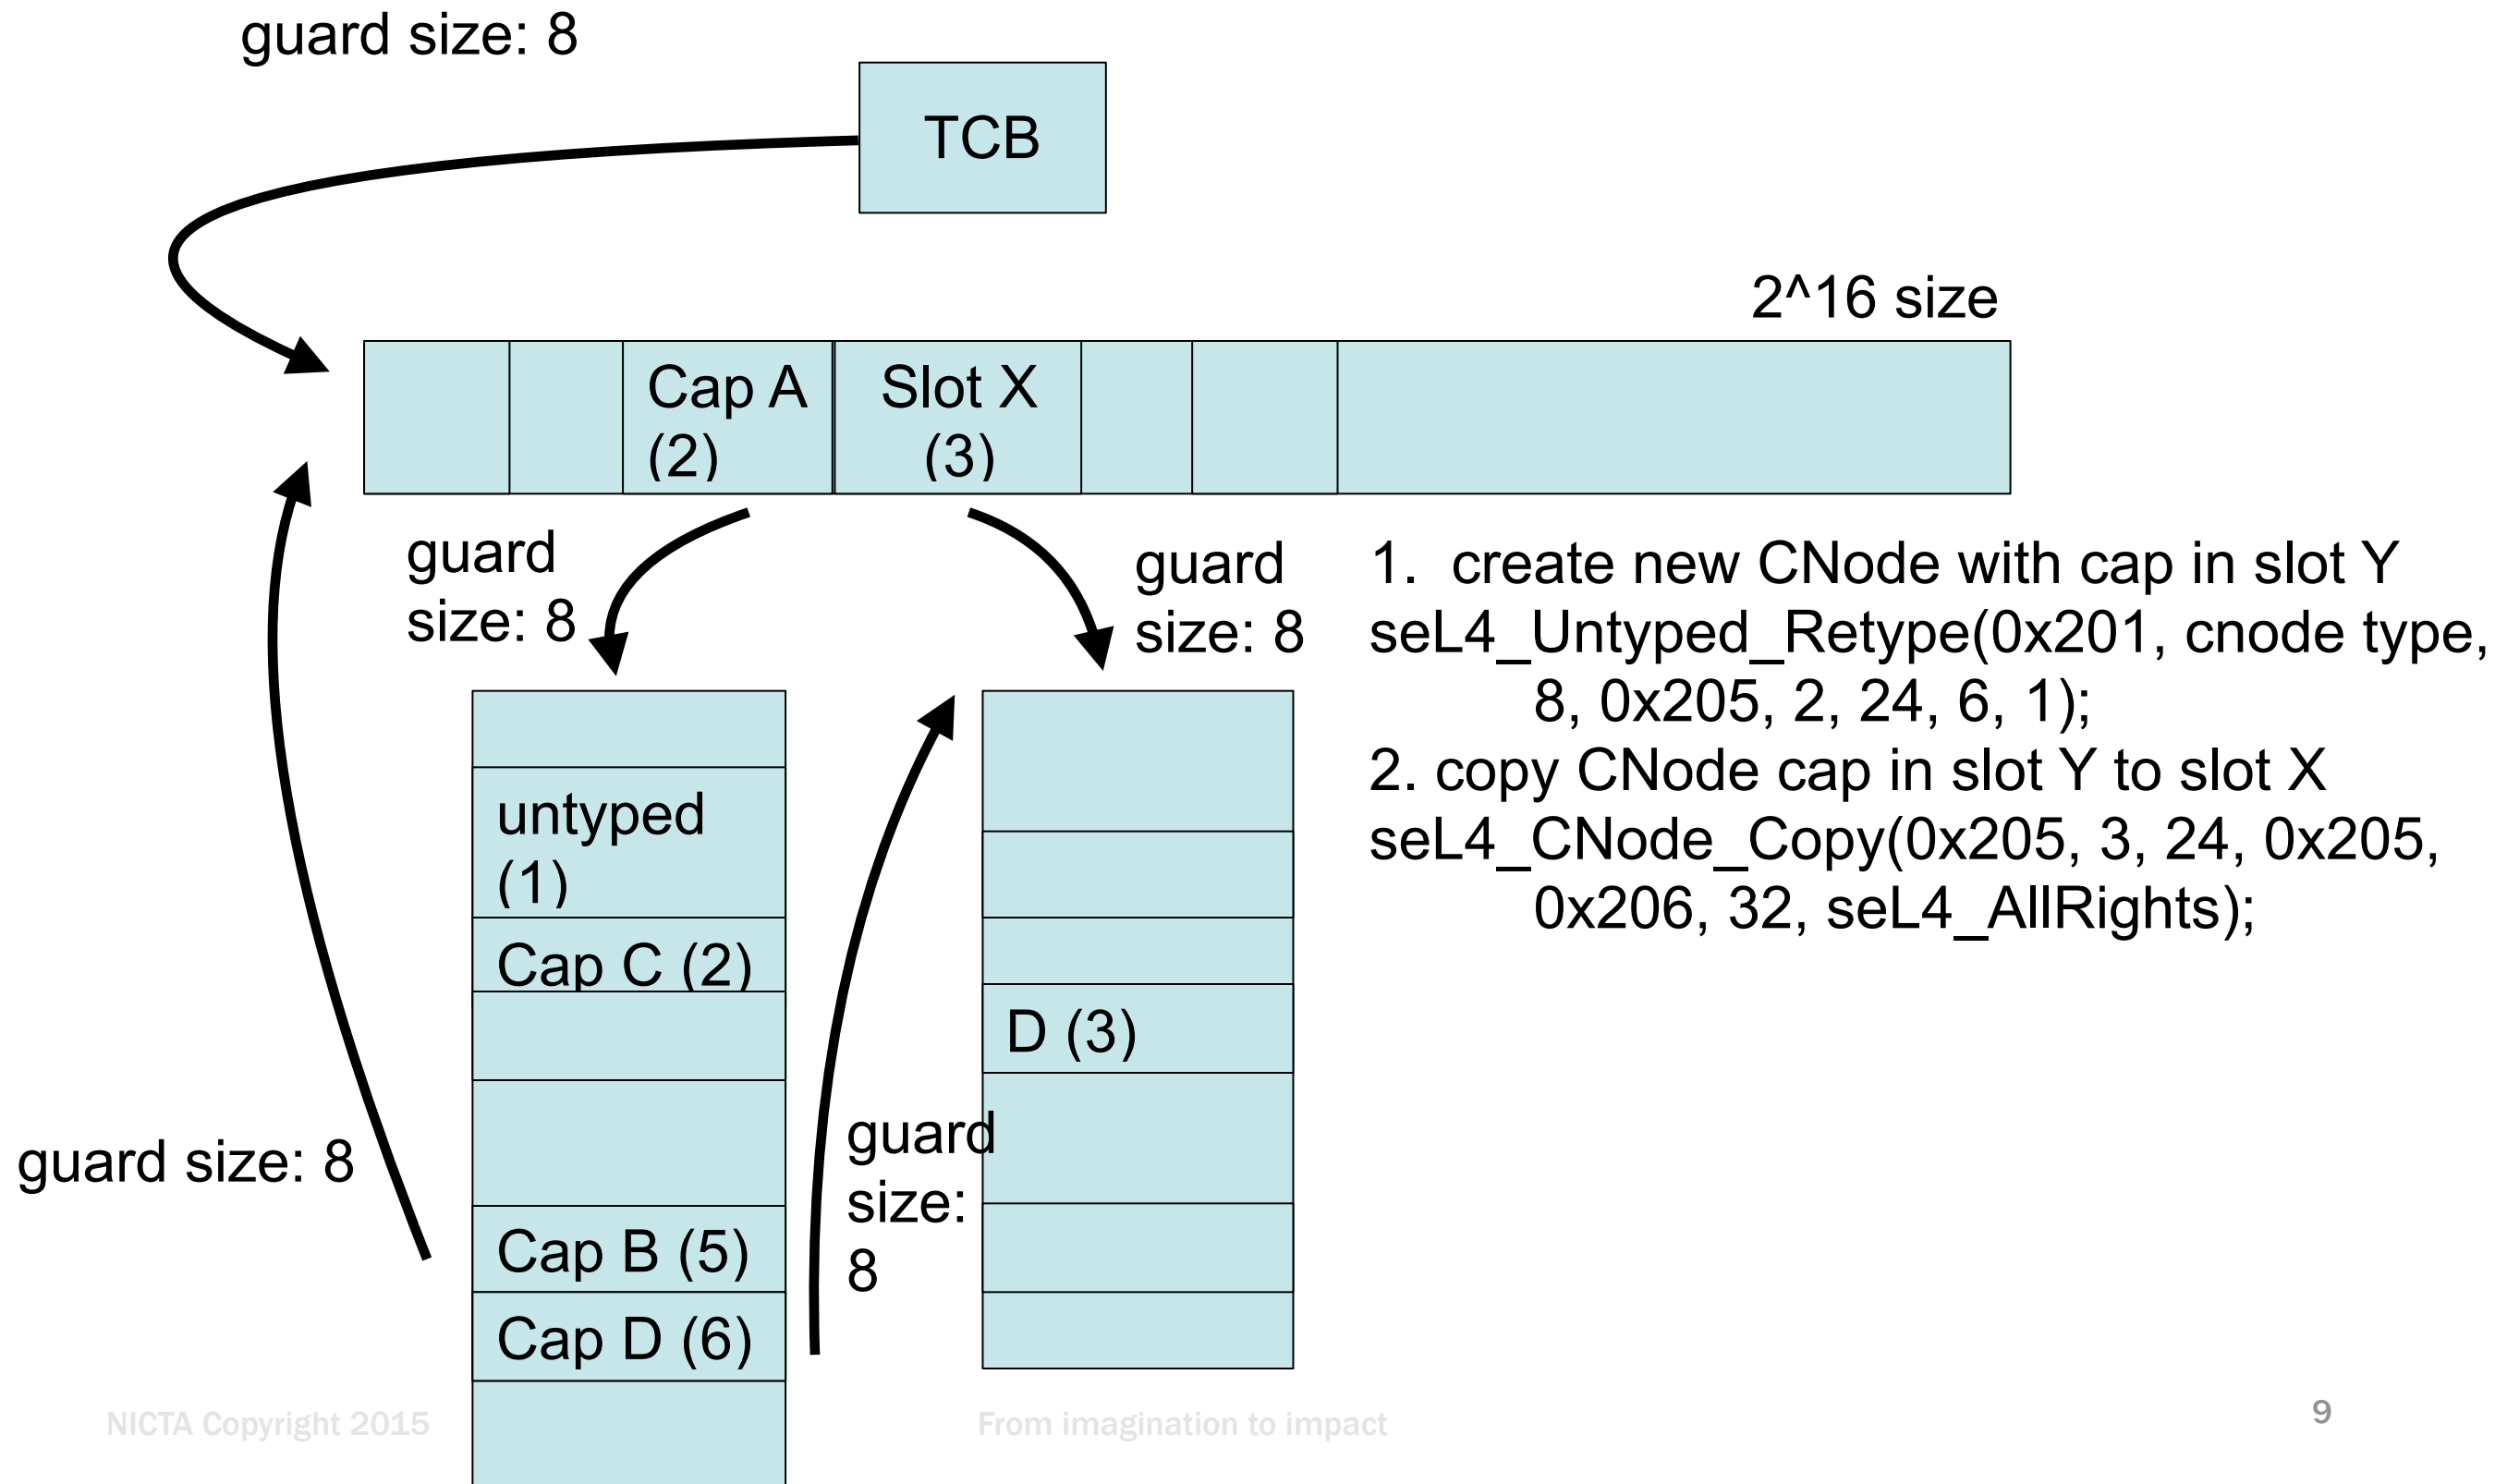
guard size: 8



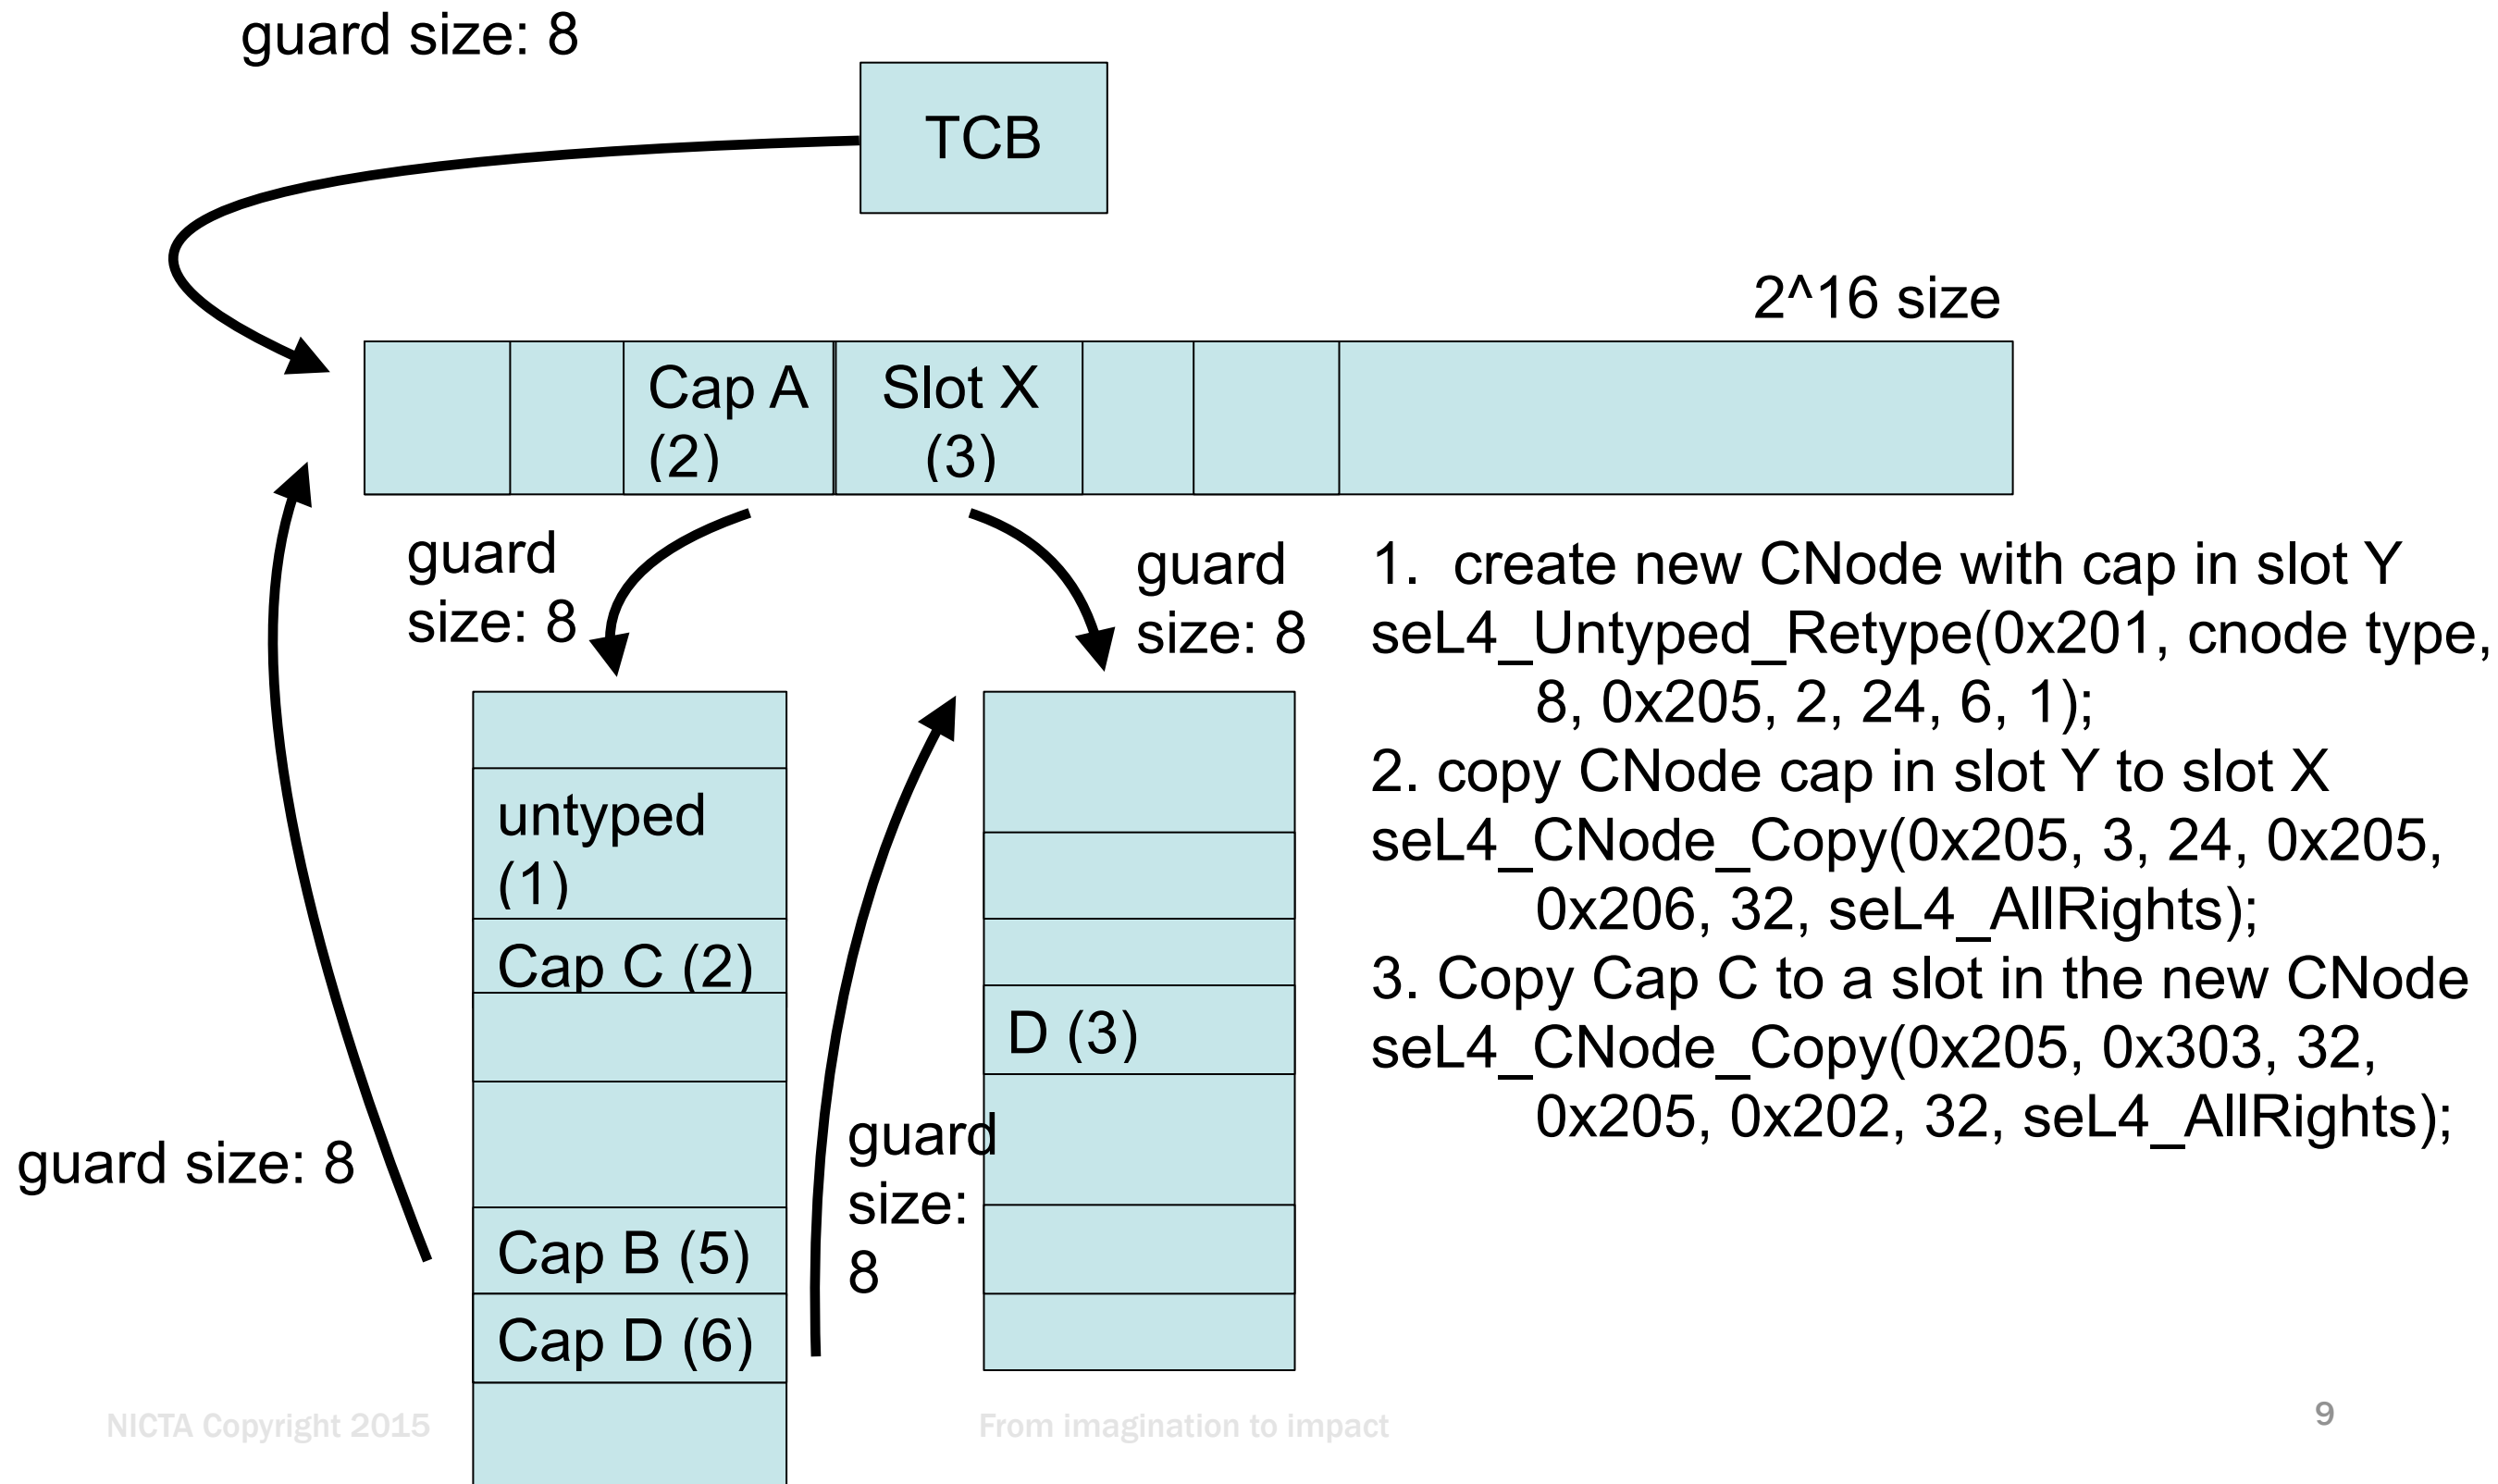
1. create new CNode with cap in slot Y
`seL4_Untyped_Retype(0x201, cnode type, 8, 0x205, 2, 24, 6, 1);`
2. copy CNode cap in slot Y to slot X
`seL4_CNode_Copy(0x205, 3, 24, 0x205, 0x206, 32, seL4_AllRights);`

guard size: 8

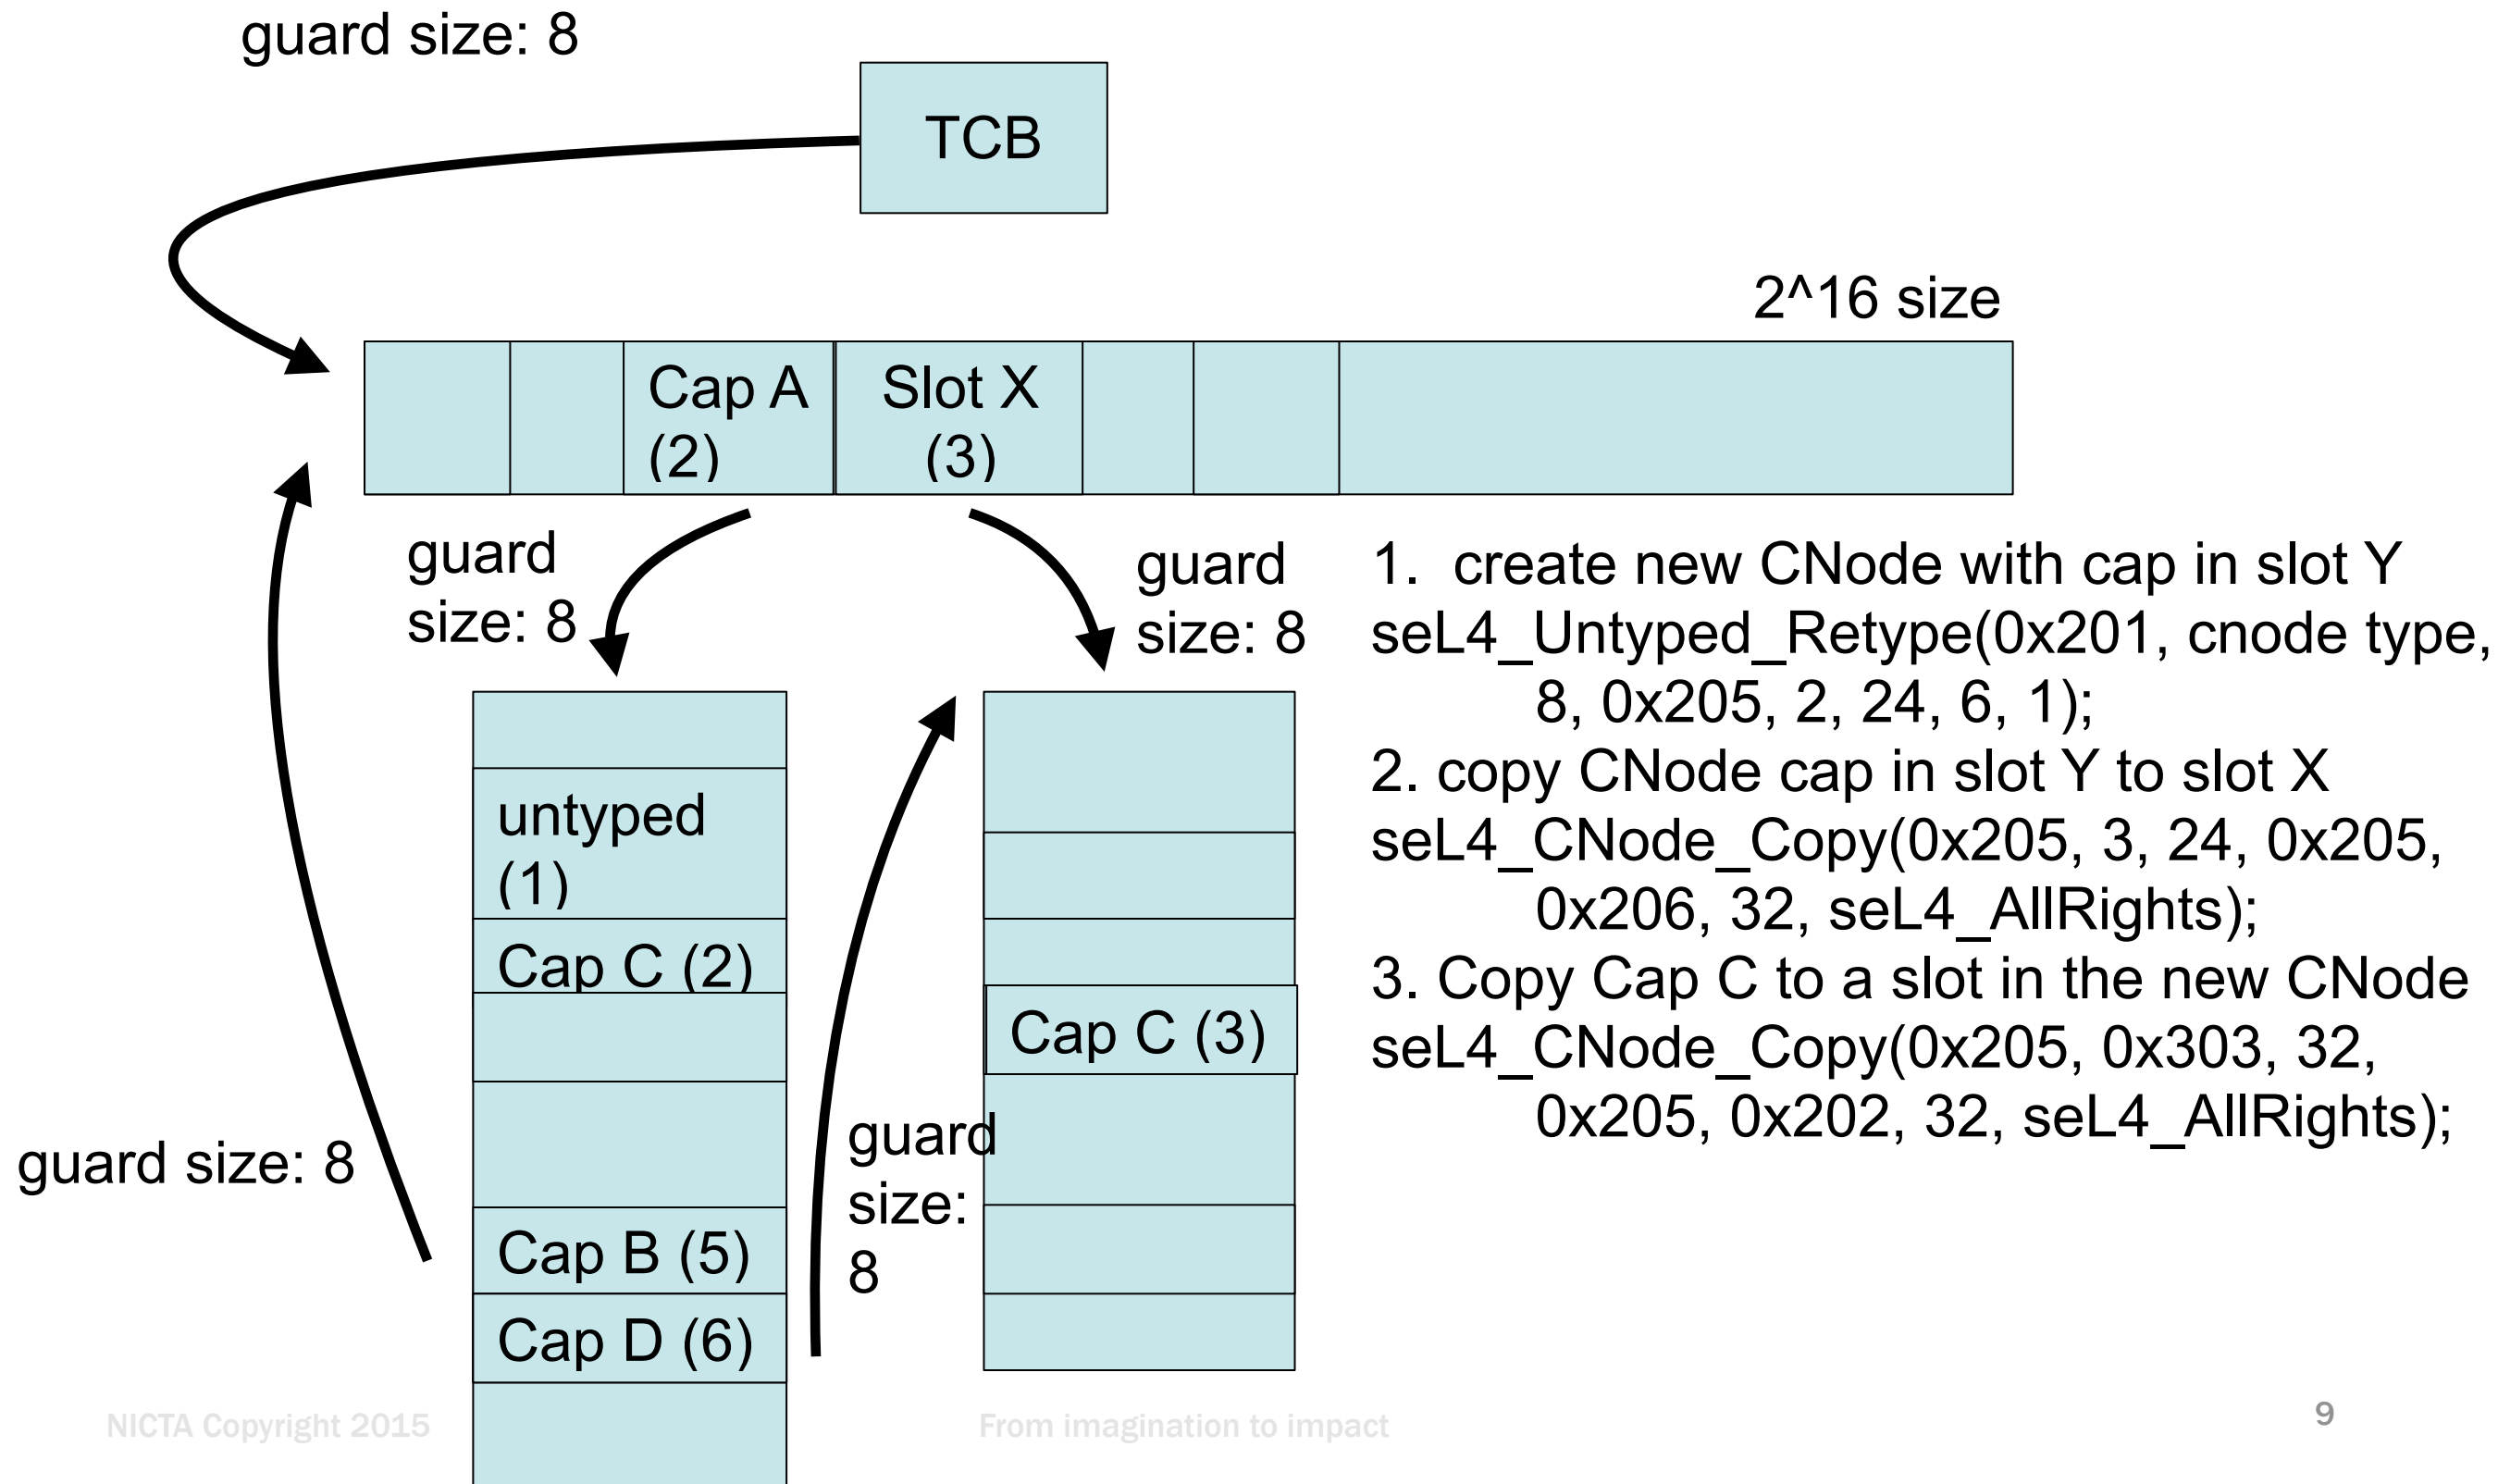
CSpace Construction Example



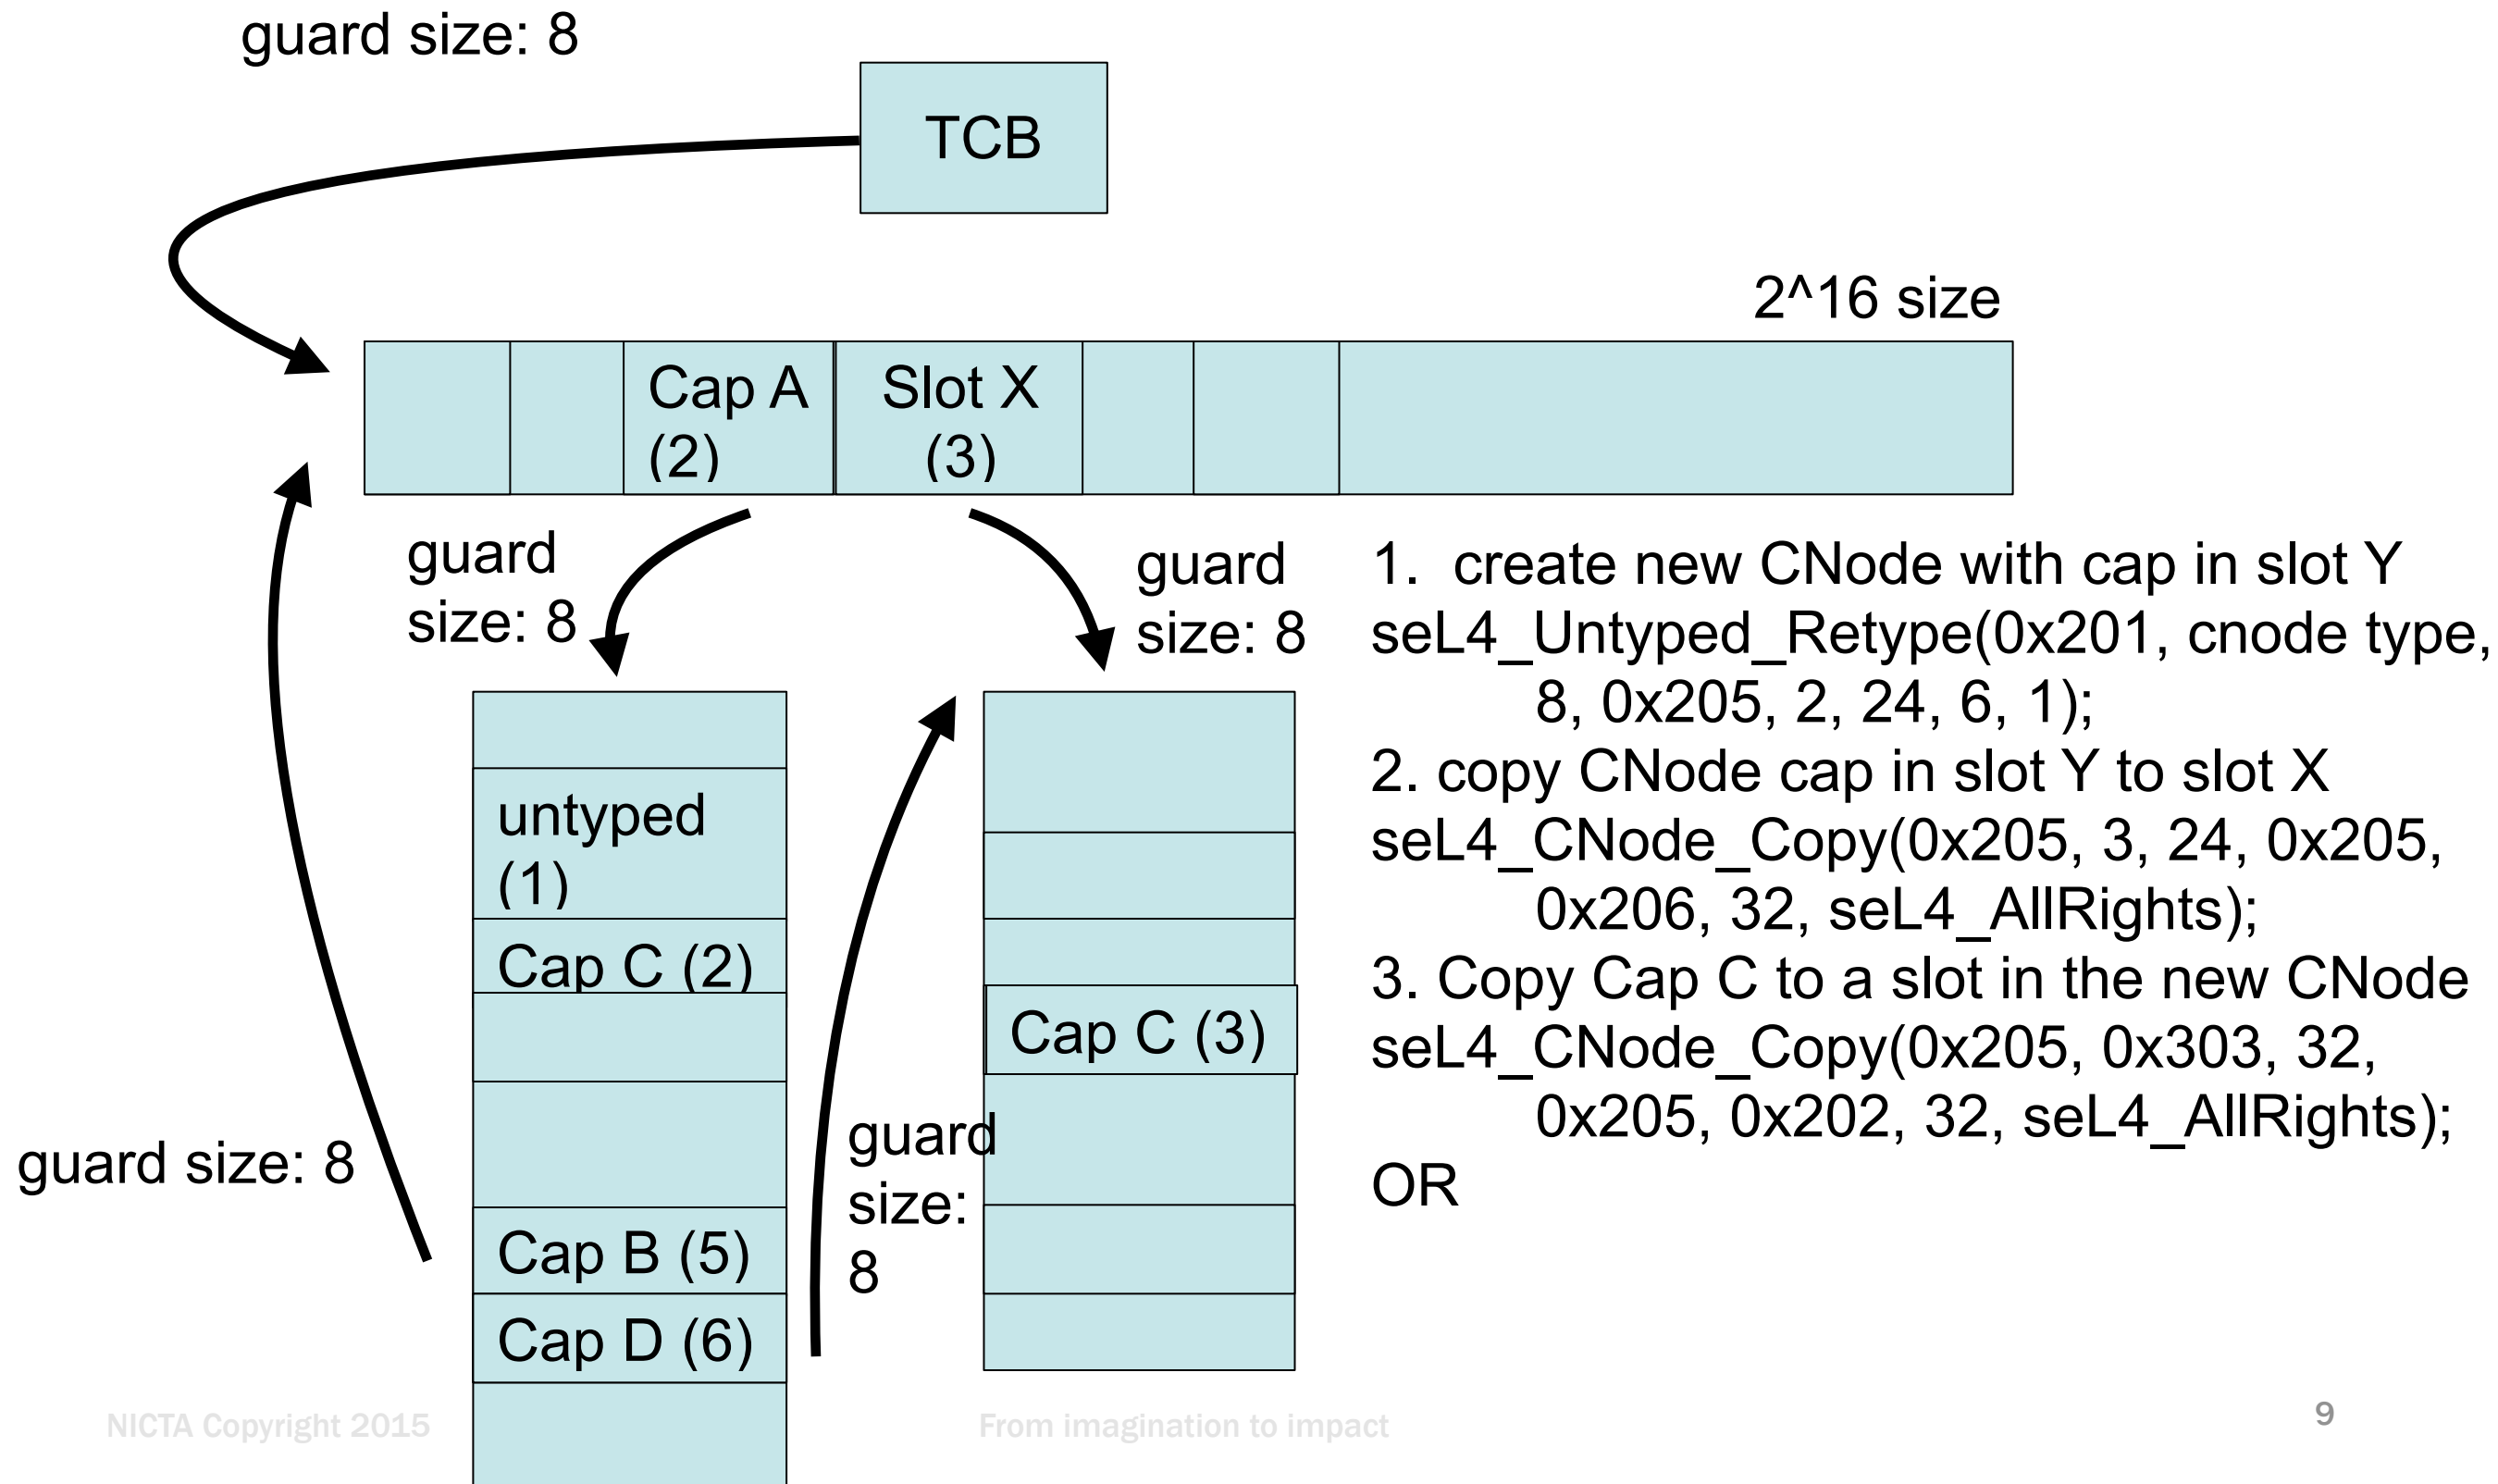
CSpace Construction Example




CSpace Construction Example



CSpace Construction Example



CSpace Construction Example



1. create new CNode with cap in slot Y
`seL4_Untyped_Retype(0x201, cnode type, 8, 0x205, 2, 24, 6, 1);`
 2. copy CNode cap in slot Y to slot X
`seL4_CNode_Copy(0x205, 3, 24, 0x205, 0x206, 32, seL4_AllRights);`
 3. Copy Cap C to a slot in the new CNode
`seL4_CNode_Copy(0x205, 0x303, 32, 0x205, 0x202, 32, seL4_AllRights);`
- OR
- `seL4_CNode_Copy(0x206, 3, 8, 0x205, 0x202, 32, seL4_AllRights);`

CSpace Construction Example

guard size: 8

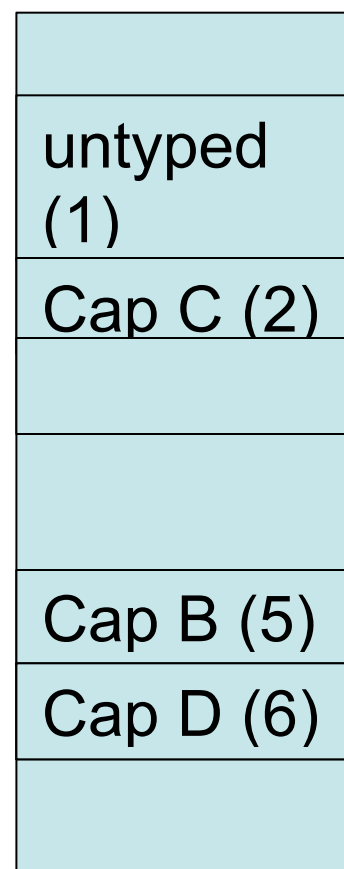
TCB

2^{16} size




guard
size: 8

guard
size: 8



guard
size:
8



1. create new CNode with cap in slot Y
`seL4_Untyped_Retype(0x201, cnode type, 8, 0x205, 2, 24, 6, 1);`
 2. copy CNode cap in slot Y to slot X
`seL4_CNode_Copy(0x205, 3, 24, 0x205, 0x206, 32, seL4_AllRights);`
 3. Copy Cap C to a slot in the new CNode
`seL4_CNode_Copy(0x205, 0x303, 32, 0x205, 0x202, 32, seL4_AllRights);`
- OR
- `seL4_CNode_Copy(0x206, 3, 8, 0x205, 0x202, 32, seL4_AllRights);`

guard size: 8

VSpace: PageDir, PageTable, Frame



- **VSpace:**
 - represents mapping: virtual address \rightarrow physical address
 - i.e. abstraction of CPU page table
 - VSpace-related objects are platform-specific
- **Sizes (ARM)**
 - PD: 16KiB, 4 byte slots
 - PT: 1KiB, 4 byte slots
 - Frame: 4KiB, 64KiB, 1MiB, 16MiB
- **Size (x86)**
 - PD: 4KiB, 4 byte slots
 - PT: 4KiB, 4 byte slots
 - Frame: 4KiB, 4MiB

VSpace API



- PD

- seL4_ARCH_PageTable_Map
- seL4_ARCH_PageTable_Unmap
- seL4_ARCH_Page_Map - for large frames
- seL4_ARCH_Page_Unmap - for large frames

- PT

- seL4_ARCH_Page_Map
- seL4_ARCH_Page_Unmap

- Note:

- ARCH is either ARM or IA32

TCB: Thread Control Block



- **TCB Object:**
 - kernel's representation of a thread.
 - contains:
 - Caps: CSpace, VSpace, IPC Buffer Frame
 - Other: IP (instruction pointer), SP (stack pointer), IPC Buffer, Priority
- **IPC Buffer**
 - buffer used to pass data during IPC
 - 512 byte object, must be wholly in one frame
 - also used for all syscalls
 - passed to TCB as:
 - index to Frame cap in TCB's CSpace
 - address where it is mapped in TCB's VSpace

TCB: API



- **Configure**
 - seL4_TCB_Configure
- **Write Registers**
 - seL4_TCB_ReadRegisters
 - seL4_TCB_WriteRegisters
- **Resume**
 - seL4_TCB_Resume
- **Suspend**
 - seL4_TCB_Suspend

Inter-Process Communication



- **Synchronous Endpoint Object**
 - enables synchronous (blocking) communication
 - communicating threads must hold caps to same endpoint
- **Endpoint Caps**
 - master cap (receiver), derived caps (senders)
 - **badge**: identifies specific sender cap
 - **reply cap**: temporary cap allows receiver to reply to sender for two-way communication

Endpoint API



- Sending and Receiving
 - seL4_Send
 - seL4_Wait
 - seL4_Call
 - seL4_ReplyWait
 - seL4_NBSend
- Message Registers
 - seL4_GetMR
 - seL4_SetMR
 - seL4_GetCap
 - seL4_SetCap

Notification



- **Asynchronous Endpoint Object**
 - allows one thread to send a notification to another
 - notification: asynchronous (non-blocking) message
 - sends limited data (32-bit word)
- **AsyncEndpoint Caps**
 - master cap (receiver), derived caps (senders)
- **API**
 - seL4_Notify
 - seL4_Wait

BootInfo: Start-up Information



- On startup, kernel creates:
 - root task CSpace
 - root task VSpace
 - frames for device memory
 - untyped caps for RAM memory
- All startup objects are available to root task
 - kernel places caps to these objects in root task CSpace
 - kernel needs to tell root task
 - what caps it has
 - what the objects are
- Bootinfo
 - info about all the initial objects and the caps to them

Initial Caps



- **Some Initial Caps**

- `seL4_CapNull = 0, /* null cap */`
- `seL4_CapInitThreadTCB = 1, /* initial thread's TCB cap */`
- `seL4_CapInitThreadCNode = 2, /* root CNode cap */`
- `seL4_CapInitThreadVSpace = 3, /* VSpace cap */`
- `seL4_CapBootInfoFrame = 9, /* bootinfo frame cap */`
- `seL4_CapInitThreadIPCBuffer = 10, /* initial thread's IPC
buffer frame cap */`

Bootinfo contents



- **seL4_Bootinfo struct**

- seL4_IPCBuffer* ipcBuffer; /* pointer to initial thread's IPC buffer */
- seL4_SlotRegion empty; /* empty slots (null caps) */
- seL4_SlotRegion userImageFrames; /* userland-image frame caps */
- seL4_SlotRegion userImagePDs; /* userland-image PD caps */
- seL4_SlotRegion userImagePTs; /* userland-image PT caps */
- seL4_SlotRegion untyped; /* untyped-object caps (untyped caps) */
- seL4_Word untypedPaddrList /* physical address of each untyped cap */
- seL4_Uint8 untypedSizeBitsList; /* size (2^n) bytes of each untyped cap */
- seL4_Uint8 initThreadCNodeSizeBits; /* root CNode size (2^n slots) */
- seL4_Word numDeviceRegions; /* number of device regions */
- seL4_DeviceRegion deviceRegions /* device regions */

- **seL4_GetBootInfo**